

CPSC 4770/6770

Distributed and Cluster Computing

Lecture 13: Introduction to Hadoop MapReduce

Key Steps to Use Hadoop on Palmetto

- Updating .bashrc on your Palmetto home directory
 - `echo "module load openjdk/1.8.0_222-b10-gcc/8.3.1 hadoop/3.2.1-gcc/8.3.1" >> ~/.bashrc`
- Requesting a compute node (e.g. node0001 ~ nodeXXXX) via qsub, e.g.:
 - `qsub -l -l select=3:ncpus=8:mem=14gb,walltime=03:00:00`
- Copying the myhadoop template from /zfs/citi
 - `cp -R /zfs/citi/myhadoop/ ~/`
 - `cd ~/myhadoop`
- Examining the myhadoop template
 - `ls -l`
 - `ls -l bin/`
- Launching myhadoop
 - `./init_hadoop.sh`
- Testing myhadoop
 - `./test_hadoop.sh`
- Hadoop main commands
 - `hdfs`
 - `hdfs dfs`
- Specifying configuration location
 - `export HADOOP_CONF_DIR="/home/$USER/hadoop_palmetto/config/"`
 - `hdfs dfs -mkdir /user/`
 - `hdfs dfs -mkdir /user/$USER`
 - `hdfs dfs -ls /user/`
 - `hdfs dfs -ls /user/$USER`
- Stopping myhadoop
 - `./stop_hadoop.sh`

Home directory on HDFS

- In HDFS, the home directory is defaulted to be /user/\$USER with \$USER as your username
- `hdfs dfs -ls /user/$USER`
- `hdfs dfs -ls`
- `hdfs dfs -ls .`

List of HDFS Commands and Usages

- <https://hadoop.apache.org/docs/r2.6.4/hadoop-project-dist/hadoop-hdfs/HDFSCommands.html>
- <http://fibrevillage.com/storage/630-using-hdfs-command-line-to-manage-files-and-directories-on-hadoop>

hdfs

hdfs dfs

hdfs dfs -help

Create a Directory in HDFS

- Usage: `hdfs dfs -mkdir <path>`
- Example: Create a directory named **intro-to-hadoop** inside your user directory on HDFS
 - `hdfs dfs -ls /user/$USER/intro-to-hadoop`
 - `hdfs dfs -ls /user/$USER`

Upload/download a File to HDFS

- Copy single file, or multiple files from local file system to HDFS
 - Usage: `hdfs dfs -put <localsrc> ... <HDFS_dest_Path>`
- Download single file, or multiple files from HDFS to local file system
 - Usage: `hdfs dfs -get <HDFS_src_Path> ... <localdest>`
- Example: Upload a text file into the newly created *intro-to-hadoop* directory and copy it back
 - `hdfs dfs -put /zfs/citi/complete-shakespeare.txt intro-to-hadoop/`
 - `hdfs dfs -ls intro-to-hadoop`
 - `hdfs dfs -head intro-to-hadoop/complete-shakespeare.txt`
 - `hdfs dfs -get intro-to-hadoop/complete-shakespeare.txt ~/shakespeare-complete.txt`
 - `head ~/shakespeare-complete.txt`
 - `diff /zfs/citi/complete-shakespeare.txt ~/shakespeare-complete.txt`

Upload/download a directory to HDFS

- The `put` and `get` subsubcommands can also be used to move directories as well
- Example: upload a movielens directory into the intro-to-Hadoop directory
 - `hdfs dfs -put /zfs/citi/movielens intro-to-hadoop/`
 - `hdfs dfs -ls intro-to-hadoop`
 - `hdfs dfs -ls intro-to-hadoop/movielens`

Check a File State in HDFS

- Check a file/directory state in HDFS
- Usage: `hdfs fsck <path>`
- Example: Check the health status of the directories */intro-to-hadoop* in HDFS using **fsck**
 - `hdfs fsck intro-to-hadoop/-files -blocks -locations`

```
Status: HEALTHY
Number of data-nodes: 2
Number of racks: 1
Total dirs: 2
Total symlinks: 0

Replicated Blocks:
Total size: 1029304096 B
Total files: 8
Total blocks (validated): 21 (avg. block size 49014480 B)
Minimally replicated blocks: 21 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 21 (100.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 3
Average block replication: 2.0
Missing blocks: 0
Corrupt blocks: 0
Missing replicas: 21 (33.333332 %)

Erasure Coded Block Groups:
Total size: 0 B
Total files: 0
Total block groups (validated): 0
Minimally erasure-coded block groups: 0
Over-erasure-coded block groups: 0
Under-erasure-coded block groups: 0
Unsatisfactory placement block groups: 0
Average block group size: 0.0
Missing block groups: 0
Corrupt block groups: 0
Missing internal blocks: 0
FSCK ended at Fri Sep 18 09:09:07 EDT 2020 in 29 milliseconds

The filesystem under path '/user/jin6/intro-to-hadoop' is HEALTHY
```


See Contents of a File in HDFS

- Usage: `hdfs dfs -cat <path[filename]>`
- Example: display the content of *complete-shakespeare.txt* in HDFS
 - `hdfs dfs -cat intro-to-hadoop/complete-shakespeare.txt`

Top 10 Hadoop HDFS Commands



Source: <https://data-flair.training/blogs/top-hadoop-hdfs-commands-tutorial/>

Reality of Working with Big Data

- Hundreds or thousands of machines to support big data
 - Distribute data for storage (HDFS)
 - Parallelize data computation (Hadoop MapReduce)
 - Handle failure (HDFS and Hadoop MapReduce)

MapReduce Programming Paradigm

- What is “map”?
 - A function/procedure that is applied to every individual elements of a collection/list/array/...
 - e.g.,

```
int square(x) { return x*x;}  
map square [1,2,3,4] -> [1,4,9,16]
```
- What is “reduce”?
 - A function/procedure that performs an operation on a list. This operation will “fold/reduce” this list into a single value (or a smaller subset)
 - e.g.,

```
reduce ([1,2,3,4]) using sum -> 10  
reduce ([1,2,3,4]) using multiply -> 24
```
- In HDFS, **map** tasks are performed on top of individual data blocks; **reduce** tasks are performed on intermediate results from map tasks to calculate the final results.

Implementation of MapReduce Programming Paradigm in Hadoop MapReduce

- **Programmers implement:**

- Map function: Take in the input data and return a *key,value* pair
- Reduce function: Receive the *key,value* pairs from the mapper and provide a final output as a reduction operation on the pairs
- Optional functions:
 - Partition function: determines the distribution of mappers' *key,value* pairs to the reducers
 - Combine functions: initial reduction on the mappers to reduce network traffics

- **The MapReduce Framework handles everything else**

What is “everything else”?

- **What is "everything else"?**
 - Scheduling
 - Data distribution
 - Synchronization
 - Error and Fault Handling
- **The cost of "everything else"?**
 - All algorithms must be expressed as a combination of mapping, reducing, combining, and partitioning functions
 - No control over execution placement of mappers and reducers
 - No control over life cycle of individual mappers and reducers
 - Very limited information about which mapper handles which data block
 - Very limited information about which reducer handles which intermediate key
- **Additional challenge**
 - Large scale debugging on big data programming is difficult

Applications of MapReduce

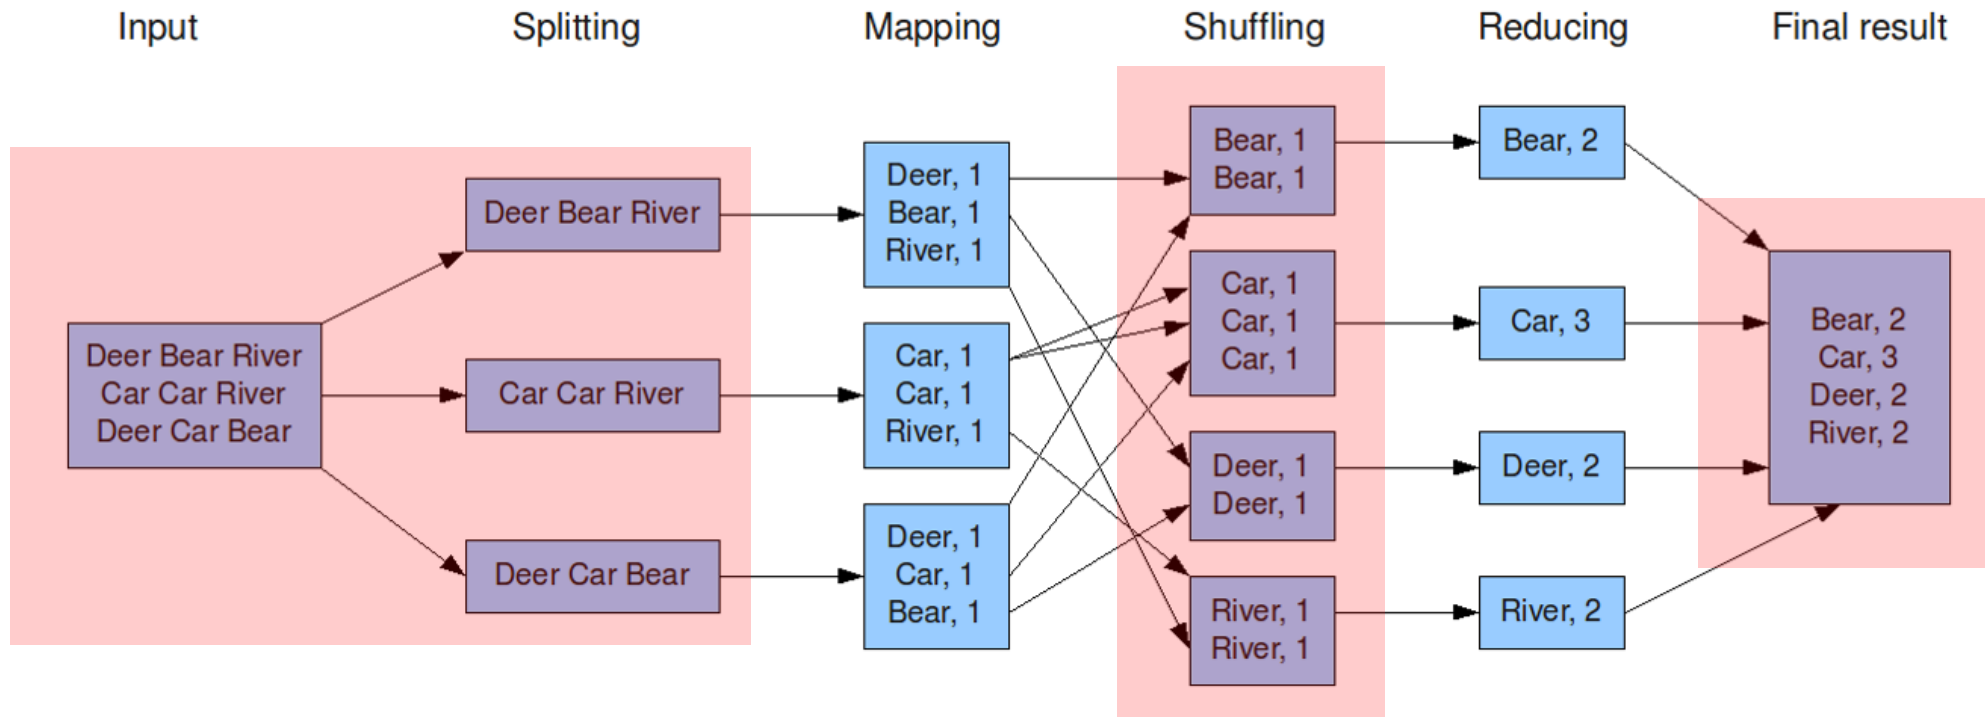
- Text tokenization, indexing, and search
 - Web access log stats
 - Inverted index construction
 - Term-vector per host
 - Distributed grep/sort
- Graph creation
 - Web link-graph reversal (Google's PageRank)
- Data Mining and machine learning
 - Document clustering
 - Machine learning
 - Statistical machine translation

WordCount: the *Hello, World* of Big Data

- Count how many unique words there are in a file/multiple files
- Standard parallel programming approach:
 - Count number of files
 - Set number of processes
 - Possibly setting up dynamic workload assignment
 - A lot of data transfer
 - Significant coding effort

MapReduce WordCount Example

The overall MapReduce word count process



WordCount (wcMapper.py)

- `hdfs dfs -cat intro-to-hadoop/complete-shakespeare.txt 2>/dev/null | head -n 20`
- `cd ~/myhadoop`
- `cat -n codes/wcMapper.py`

```
#!/usr/bin/env python
import sys
for oneLine in sys.stdin:
    oneLine = oneLine.strip()
    for word in oneLine.split(" "):
        if word != "":
            print ('%s\t%s' % (word, 1))
```

- module load anaconda3/2020.07-gcc/8.3.1
- `hdfs dfs -cat intro-to-hadoop/ complete-shakespeare.txt 2>/dev/null | head -n 20 | python ./codes/wcMapper.py`
- `hdfs dfs -cat intro-to-hadoop/ complete-shakespeare.txt 2>/dev/null | head -n 20 | python ./codes/wcMapper.py | sort`

WordCount (wctReducer.py)

```
#!/usr/bin/env python
import sys

current_word = None
total_word_count = 0

for line in sys.stdin:
    line = line.strip()
    word, count = line.split("\t", 1)
    try:
        count = int(count)
    except ValueError:
        continue

    if current_word == word:
        total_word_count += count
    else:
        if current_word:
            print ("%s\t%s" % (current_word, total_word_count))
            current_word = word
            total_word_count = 1

if current_word == word:
    print ("%s\t%s" % (current_word, total_word_count))
```

- `cat -n codes/wcReducer.py`
- `hdfs dfs -cat intro-to-hadoop/complete-shakespeare.txt 2>/dev/null | head -n 20 | python wcMapper.py | sort | python ./codes/wcReducer.py`
- Or you can write in a more readable format:
`hdfs dfs -cat intro-to-hadoop/ complete-shakespeare.txt \`
`2>/dev/null \`
`| head -n 20 \`
`| python wcMapper.py \`
`| sort \`
`| python ./codes/wcReducer.py`

Run a MapReduce Program by Streaming

- `hdfs dfs -rm -R intro-to-hadoop/output-wordcount`
- `mapred --config ~/hadoop_palmetto/config streaming \`
 `-input intro-to-hadoop/text/gutenberg-shakespeare.txt \`
 `-output intro-to-hadoop/output-wordcount \`
 `-file wordcountMapper.py \`
 `-mapper wordcountMapper.py \`
 `-file wordcountReducer.py \`
 `-reducer wordcountReducer.py`
- `hdfs dfs -cat intro-to-hadoop/output-wordcount/part-00000 2>/dev/null | head -n 100`

Challenge

- Modify *wcMapper.py* so that punctuations and capitalization are no longer factors in determining unique words

```
#!/usr/bin/env python
import sys
for oneLine in sys.stdin:
    oneLine = oneLine.strip()
    for word in oneLine.split(" "):
        if word != "":
            print ('%s\t%s' % (word, 1))
```

Partial Solution (wcEnhancedMapper.py)

```
#!/usr/bin/env python
import sys
import string

translator = str.maketrans('', '', string.punctuation)

for oneLine in sys.stdin:
    oneLine = oneLine.strip()
    for word in oneLine.split(" "):
        if word != "":
            newWord = word.translate(translator).lower()
            print ('%s\t%s' % (_____, 1))
```

hdfs dfs -cat intro-to-hadoop/complete-shakespeare.txt 2>/dev/null | head -n 20 | python wcEnhancedMapper.py | sort | python wcReducer.py

Basic Anatomy of a Java Hadoop MapReduce Application (Example: WordCount.java)

```
1 import java.io.IOException;
2 import java.util.*;
3
4 import org.apache.hadoop.conf.*;
5 import org.apache.hadoop.fs.*;
6 import org.apache.hadoop.conf.*;
7 import org.apache.hadoop.io.*;
8 import org.apache.hadoop.mapreduce.*;
9 import org.apache.hadoop.mapreduce.lib.input.*;
10 import org.apache.hadoop.mapreduce.lib.output.*;
11 import org.apache.hadoop.util.*;
```

```
13 /* MAIN */
14 public class WordCount extends Configured implements Tool {
15
16     public static void main(String args[]) throws Exception {
17         int res = ToolRunner.run(new WordCount(), args);
18         System.exit(res);
19     }
20
21     public int run(String[] args) throws Exception {
22         Path inputPath = new Path(args[0]);
23         Path outputPath = new Path(args[1]);
24
25         Configuration conf = getConf();
26         Job job = new Job(conf, this.getClass().toString());
27
28         FileInputFormat.setInputPaths(job, inputPath);
29         FileOutputFormat.setOutputPath(job, outputPath);
30
31         job.setJarByClass(WordCount.class);
32         job.setInputFormatClass(TextInputFormat.class);
33         job.setOutputFormatClass(TextOutputFormat.class);
34         job.setMapOutputKeyClass(Text.class);
35         job.setMapOutputValueClass(IntWritable.class);
36         job.setOutputKeyClass(Text.class);
37         job.setOutputValueClass(IntWritable.class);
38
39         job.setMapperClass(Map.class);
40         job.setCombinerClass(Reduce.class);
41         job.setReducerClass(Reduce.class);
42
43         return job.waitForCompletion(true) ? 0 : 1;
44     }
```

WordCount.java (Map)

```
46  /* MAP */
47  public static class Map extends Mapper<Object, Text, Text, IntWritable> {
48      private Text word = new Text();
49
50      public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
51          StringTokenizer tokenizer = new StringTokenizer(value.toString());
52          while (tokenizer.hasMoreTokens()) {
53              word.set(tokenizer.nextToken());
54              context.write(word, new IntWritable(1));
55          }
56      }
```


WordCount.java (Reduce)

```
58  /* REDUCE */
59  public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable>{
60      public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
61          int sum = 0;
62          for (IntWritable value : values) {
63              sum += value.get();
64          }
65          context.write(key, new IntWritable(sum));
66      }
67  }
68 }
```