

CPSC 4770/6770

Distributed and Cluster Computing

Lecture 15: Optimization

First Principle

- First principle of optimizing Hadoop workflow: **Reduce data movement in the shuffle phase**

```
hdfs dfs -rm -r intro-to-hadoop/output-movieLens-02  
mapred streaming \  
-input intro-to-hadoop/movieLens/ratings.csv \  
-output intro-to-hadoop/output-movieLens-02 \  
-file ./codes/avgRatingMapper04.py \  
-mapper avgRatingMapper04.py \  
-file ./codes/avgRatingReducer01.py \  
-reducer avgRatingReducer01.py \  
-file ./movieLens/movies.csv
```

- What is being passed from Map to Reduce?
- Can reducer do the same thing as mapper, that is, to load in external data?
- If we load external data on the reduce side, do we need to do so on the map side?

avgRatingReducer02.py

```
%%writefile codes/avgRatingReducer02.py
#!/usr/bin/env python
import sys
import csv

movieFile = "./movies.csv"
movieList = {}

with open(movieFile, mode = 'r') as infile:
    reader = csv.reader(infile)
    for row in reader:
        movieList[row[0]] = {}
        movieList[row[0]]["title"] = row[1]
        movieList[row[0]]["genre"] = row[2]

current_movie = None
current_rating_sum = 0
current_rating_count = 0

for line in sys.stdin:
    line = line.strip()
    movie, rating = line.split("\t", 1)
    try:
        rating = float(rating)
    except ValueError:
        continue

    if current_movie == movie:
        current_rating_sum += rating
        current_rating_count += 1
    else:
        if current_movie:
            rating_average = current_rating_sum / current_rating_count
            movieTitle = movieList[current_movie]["title"]
            movieGenres = movieList[current_movie]["genre"]
            print ("%s\t%s\t%s" % (movieTitle, rating_average, movieGenres))
            current_movie = movie
            current_rating_sum = rating
            current_rating_count = 1

if current_movie == movie:
    rating_average = current_rating_sum / current_rating_count
    movieTitle = movieList[current_movie]["title"]
    movieGenres = movieList[current_movie]["genre"]
    print ("%s\t%s\t%s" % (movieTitle, rating_average, movieGenres))
```

hdfs dfs -rm -r intro-to-hadoop/output-movielens-03
mapred streaming \

-input /repository/movielens/ratings.csv \
-output intro-to-hadoop/output-movielens-03 \
-file avgRatingMapper02.py \
-mapper avgRatingMapper02.py \
-file avgRatingReducer02.py \
-reducer avgRatingReducer02.py \
-file ./movielens/movies.csv

hdfs dfs -ls intro-to-hadoop/output-movielens-02
hdfs dfs -ls intro-to-hadoop/output-movielens-03

hdfs dfs -cat intro-to-hadoop/output-movielens-03/part-00000 \
2>/dev/null | head -n 10

How does the number shuffle bytes in this example compare to the previous example?

Performance Comparison

```
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Data-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=46601696
  Total time spent by all reduces in occupied slots (ms)=1731216
  Total time spent by all map tasks (ms)=1456303
  Total time spent by all reduce tasks (ms)=108201
  Total vcore-milliseconds taken by all map tasks=1456303
  Total vcore-milliseconds taken by all reduce tasks=108201
  Total megabyte-milliseconds taken by all map tasks=5965017088
  Total megabyte-milliseconds taken by all reduce tasks=221595648

Map-Reduce Framework
  Map input records=24404097
  Map output records=24404096
  Map output bytes=1220481613
  Map output materialized bytes=1269290283
  Input split bytes=1410
  Combine input records=0
  Combine output records=0
  Reduce input groups=39408
  Reduce shuffle bytes=1269290283
  Reduce input records=24404096
  Reduce output records=39408
  Spilled Records=73212288
  Shuffled Maps =10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=20688
  CPU time spent (ms)=591720
  Physical memory (bytes) snapshot=3584802816
  Virtual memory (bytes) snapshot=21488558080
  Total committed heap usage (bytes)=1742340096
  Peak Map Physical memory (bytes)=491581440
  Peak Map Virtual memory (bytes)=2332282880
  Peak Reduce Physical memory (bytes)=188510208
  Peak Reduce Virtual memory (bytes)=2169491456

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=664600312
File Output Format Counters
  Bytes Written=2139648
```

Baseline run

```
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Data-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=28225312
  Total time spent by all reduces in occupied slots (ms)=1529776
  Total time spent by all map tasks (ms)=882041
  Total time spent by all reduce tasks (ms)=95611
  Total vcore-milliseconds taken by all map tasks=882041
  Total vcore-milliseconds taken by all reduce tasks=95611
  Total megabyte-milliseconds taken by all map tasks=3612839936
  Total megabyte-milliseconds taken by all reduce tasks=195811328

Map-Reduce Framework
  Map input records=24404097
  Map output records=24404096
  Map output bytes=217230069
  Map output materialized bytes=266038321
  Input split bytes=1410
  Combine input records=0
  Combine output records=0
  Reduce input groups=39443
  Reduce shuffle bytes=266038321
  Reduce input records=24404096
  Reduce output records=39443
  Spilled Records=48808192
  Shuffled Maps =10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=17704
  CPU time spent (ms)=407410
  Physical memory (bytes) snapshot=3175272448
  Virtual memory (bytes) snapshot=20763410432
  Total committed heap usage (bytes)=2051997696
  Peak Map Physical memory (bytes)=465309696
  Peak Map Virtual memory (bytes)=2159755264
  Peak Reduce Physical memory (bytes)=425574400
  Peak Reduce Virtual memory (bytes)=2124705792

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=664600312
File Output Format Counters
  Bytes Written=2141098
```

Optimized run

Common Optimization Approaches

- In-mapper reduction of key/value pairs
- Additional combiner function

Example: Find genres which have the highest average ratings over the years

```
%%writefile codes/avgGenreMapper01.py
#!/usr/bin/env python
import sys
import csv

# for nonHDFS run
movieFile = "./movielens/movies.csv"

# for HDFS run
#movieFile = "/movies.csv"
movieList = {}

with open(movieFile, mode = 'r') as infile:
    reader = csv.reader(infile)
    for row in reader:
        movieList[row[0]] = {}
        movieList[row[0]]["title"] = row[1]
        movieList[row[0]]["genre"] = row[2]

for oneMovie in sys.stdin:
    oneMovie = oneMovie.strip()
    ratingInfo = oneMovie.split(",")
    try:
        genreList = movieList[ratingInfo[1]]["genre"]
        rating = float(ratingInfo[2])
        for genre in genreList.split("|"):
            print ("%s\t%s" % (genre, rating))
    except ValueError:
        continue
```

```
%%writefile codes/avgGenreReducer01.py
#!/usr/bin/env python
import sys
import csv
import json

current_genre = None
current_rating_sum = 0
current_rating_count = 0

for line in sys.stdin:
    line = line.strip()
    genre, rating = line.split("\t", 1)

    if current_genre == genre:
        try:
            current_rating_sum += float(rating)
            current_rating_count += 1
        except ValueError:
            continue
    else:
        if current_genre:
            rating_average = current_rating_sum / current_rating_count
            print ("%s\t%s" % (current_genre, rating_average))
            current_genre = genre
        try:
            current_rating_sum = float(rating)
            current_rating_count = 1
        except ValueError:
            continue

if current_genre == genre:
    rating_average = current_rating_sum / current_rating_count
    print ("%s\t%s" % (current_genre, rating_average))
```

- `cat -n codes/avgGenreMapper01.py`
- `cat -n codes/avgGenreReducer01.py`
- `hdfs dfs -rm -r intro-to-hadoop/output-movielens-04`
- `mapred streaming -input intro-to-hadoop/movielens/ratings.csv -output intro-to-hadoop/output-movielens-04 -file ./codes/avgGenreMapper01.py -mapper avgGenreMapper01.py -file ./avgGenreReducer01.py -reducer avgGenreReducer01.py -file ./movielens/movies.csv`

Optimization through in-mapper reduction of key/value pairs

```
%%writefile codes/avgGenreMapper02.py
#!/usr/bin/env python

import sys
import csv
import json

# for nonHDFS run
# movieFile = "./movielens/movies.csv"

# for HDFS run
movieFile = "./movies.csv"

movieList = {}
genreList = {}

with open(movieFile, mode = 'r') as infile:
    reader = csv.reader(infile)
    for row in reader:
        movieList[row[0]] = {}
        movieList[row[0]]["title"] = row[1]
        movieList[row[0]]["genre"] = row[2]

for oneMovie in sys.stdin:
    oneMovie = oneMovie.strip()
    ratingInfo = oneMovie.split(",")
    try:
        genres = movieList[ratingInfo[1]]["genre"]
        rating = float(ratingInfo[2])
        for genre in genres.split("|"):
            if genre in genreList:
                genreList[genre]["total_rating"] += rating
                genreList[genre]["total_count"] += 1
            else:
                genreList[genre] = {}
                genreList[genre]["total_rating"] = rating
                genreList[genre]["total_count"] = 1
    except ValueError:
        continue

for genre in genreList:
    print ("%s\t%s" % (genre, json.dumps(genreList[genre])))
```

- `cat -n codes/avgGenreMapper01.py`
- Make sure to use nonHDFS run in the code avgGenreMapper01.py for the following command:
 - `hdfs dfs -cat intro-to-hadoop/movielens/ratings.csv 2>/dev/null \`
`| head -n 10 | python avgGenreMapper01.py`
- `hdfs dfs -cat /repository/movielens/ratings.csv 2>/dev/null | head -n 10 | python avgGenreMapper02.py`

```
[jin6@node0088 codes]$ hdfs dfs -cat /repository/movielens/ratings.csv 2>/dev/null | head -n 10 | python avgGenreMapper01.py
Comedy 2.0
Romance 2.0
Action 1.0
Sci-Fi 1.0
Thriller 1.0
Crime 5.0
Drama 5.0
Comedy 4.0
Romance 4.0
Drama 3.0
Thriller 3.0
Drama 3.0
Horror 1.0
Comedy 2.0
Action 2.0
Adventure 2.0
Thriller 2.0
[jin6@node0088 codes]$ hdfs dfs -cat /repository/movielens/ratings.csv 2>/dev/null | head -n 10 | python avgGenreMapper02.py
Horror {"total_count": 1, "total_rating": 1.0}
Action {"total_count": 2, "total_rating": 3.0}
Drama {"total_count": 3, "total_rating": 11.0}
Crime {"total_count": 1, "total_rating": 5.0}
Adventure {"total_count": 1, "total_rating": 2.0}
Comedy {"total_count": 3, "total_rating": 8.0}
Sci-Fi {"total_count": 1, "total_rating": 1.0}
Thriller {"total_count": 3, "total_rating": 6.0}
Romance {"total_count": 2, "total_rating": 6.0}
```


Optimization through in-mapper reduction of key/value pairs (cont.)

```
%%writefile codes/avgGenreReducer02.py
#!/usr/bin/env python
import sys
import csv
import json

current_genre = None
current_rating_sum = 0
current_rating_count = 0

for line in sys.stdin:
    line = line.strip()
    genre, ratingString = line.split("\t", 1)
    ratingInfo = json.loads(ratingString)

    if current_genre == genre:
        try:
            current_rating_sum += ratingInfo["total_rating"]
            current_rating_count += ratingInfo["total_count"]
        except ValueError:
            continue
    else:
        if current_genre:
            rating_average = current_rating_sum / current_rating_count
            print ("%s\t%s" % (current_genre, rating_average))
            current_genre = genre
        try:
            current_rating_sum = ratingInfo["total_rating"]
            current_rating_count = ratingInfo["total_count"]
        except ValueError:
            continue

if current_genre == genre:
    rating_average = current_rating_sum / current_rating_count
    print ("%s\t%s" % (current_genre, rating_average))
```

- `cat -n codes/avgGenreReducer02.py`
 - `hdfs dfs -cat intro-to-hadoop/movielens/ratings.csv 2>/dev/null | head -n 10 | python avgGenreMapper02.py | sort | python avgGenreReducer02.py`
- ```
[jin6@node0088 codes]$ hdfs dfs -cat /repository/movielens/ratings.csv 2>/dev/null | head -n 10 | python avgGenreMapper02.py | sort | python avgGenreReducer02.py
Action 1.5
Adventure 2.0
Comedy 2.6666666666666665
Crime 5.0
Drama 3.6666666666666665
Horror 1.0
Romance 3.0
Sci-Fi 1.0
Thriller 2.0
```
- `hdfs dfs -rm -R intro-to-hadoop/output-movielens-05`
  - `mapred streaming -input intro-to-hadoop/movielens/ratings.csv -output intro-to-hadoop/output-movielens-05 -file ./codes/avgGenreMapper02.py -mapper avgGenreMapper02.py -file ./codes/avgGenreReducer02.py -reducer avgGenreReducer02.py -file ./movielens/movies.csv`
  - `hdfs dfs -cat intro-to-hadoop/output-movielens-05/part-00000`
  - `hdfs dfs -cat intro-to-hadoop/output-movielens-04/part-00000`
  - How different are the number of shuffle bytes between the two jobs?



# Performance Comparison

```
Job Counters
 Killed map tasks=1
 Launched map tasks=11
 Launched reduce tasks=1
 Data-local map tasks=11
 Total time spent by all maps in occupied slots (ms)=54122464
 Total time spent by all reduces in occupied slots (ms)=3189280
 Total time spent by all map tasks (ms)=1691327
 Total time spent by all reduce tasks (ms)=199330
 Total vcore-milliseconds taken by all map tasks=1691327
 Total vcore-milliseconds taken by all reduce tasks=199330
 Total megabyte-milliseconds taken by all map tasks=6927675392
 Total megabyte-milliseconds taken by all reduce tasks=408227840
Map-Reduce Framework
 Map input records=24404097
 Map output records=65715698
 Map output bytes=753947479
 Map output materialized bytes=885378935
 Input split bytes=1410
 Combine input records=0
 Combine output records=0
 Reduce input groups=20
 Reduce shuffle bytes=885378935
 Reduce input records=65715698
 Reduce output records=20
 Spilled Records=197147094
 Shuffled Maps =10
 Failed Shuffles=0
 Merged Map outputs=10
 GC time elapsed (ms)=21376
 CPU time spent (ms)=898240
 Physical memory (bytes) snapshot=2785132544
 Virtual memory (bytes) snapshot=20198805504
 Total committed heap usage (bytes)=2253717504
 Peak Map Physical memory (bytes)=463224832
 Peak Map Virtual memory (bytes)=2348793856
 Peak Reduce Physical memory (bytes)=590577664
 Peak Reduce Virtual memory (bytes)=2103820288
Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
File Input Format Counters
 Bytes Read=664600312
File Output Format Counters
 Bytes Written=540
```

Baseline run

```
Job Counters
 Killed map tasks=1
 Launched map tasks=11
 Launched reduce tasks=1
 Data-local map tasks=11
 Total time spent by all maps in occupied slots (ms)=18306560
 Total time spent by all reduces in occupied slots (ms)=391744
 Total time spent by all map tasks (ms)=572080
 Total time spent by all reduce tasks (ms)=24484
 Total vcore-milliseconds taken by all map tasks=572080
 Total vcore-milliseconds taken by all reduce tasks=24484
 Total megabyte-milliseconds taken by all map tasks=2343239680
 Total megabyte-milliseconds taken by all reduce tasks=50143232
Map-Reduce Framework
 Map input records=24404097
 Map output records=200
 Map output bytes=11670
 Map output materialized bytes=12130
 Input split bytes=1410
 Combine input records=0
 Combine output records=0
 Reduce input groups=20
 Reduce shuffle bytes=12130
 Reduce input records=200
 Reduce output records=19
 Spilled Records=400
 Shuffled Maps =10
 Failed Shuffles=0
 Merged Map outputs=10
 GC time elapsed (ms)=15440
 CPU time spent (ms)=160160
 Physical memory (bytes) snapshot=2483089408
 Virtual memory (bytes) snapshot=20544786432
 Total committed heap usage (bytes)=1742798848
 Peak Map Physical memory (bytes)=341524480
 Peak Map Virtual memory (bytes)=2214707200
 Peak Reduce Physical memory (bytes)=110313472
 Peak Reduce Virtual memory (bytes)=1860853760
Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
File Input Format Counters
 Bytes Read=664600312
File Output Format Counters
 Bytes Written=514
```

Optimized run

# Optimization through combiner function

- The combiner in MapReducer is also known as 'Mini-reducer'. The primary job of Combiner is to process the output data from the mapper, before passing it to Reducer. It runs after the mapper and before the Reducer and its use is optional.
- `mapred streaming -input intro-to-hadoop/complete-shakespeare.txt -output intro-to-hadoop/output-wordcount-01 -file ./codes/wcMapper.py -mapper wcMapper.py -file ./codes/wcReducer.py -reducer wcReducer.py`
- `mapred streaming -input intro-to-hadoop/complete-shakespeare.txt -output intro-to-hadoop/output-wordcount-02 -file ./codes/wcMapper.py -mapper wcMapper.py -file ./codes/wcReducer.py -reducer wcReducer.py -combiner wcReducer.py`

# Optimization through combiner function (Cont.)

```
Map-Reduce Framework
 Map input records=124796
 Map output records=904087
 Map output bytes=6766896
 Map output materialized bytes=8575082
 Input split bytes=198
 Combine input records=0
 Combine output records=0
 Reduce input groups=67799
 Reduce shuffle bytes=8575082
 Reduce input records=904087
 Reduce output records=67799
 Spilled Records=1808174
 Shuffled Maps =2
 Failed Shuffles=0
 Merged Map outputs=2
 GC time elapsed (ms)=1299
 CPU time spent (ms)=44090
 Physical memory (bytes) snapshot=5462134784
 Virtual memory (bytes) snapshot=39903113216
 Total committed heap usage (bytes)=6152519680
Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
File Input Format Counters
 Bytes Read=5721265
File Output Format Counters
 Bytes Written=721220
```

Without combiner

```
Map-Reduce Framework
 Map input records=124796
 Map output records=904087
 Map output bytes=6766896
 Map output materialized bytes=1105621
 Input split bytes=198
 Combine input records=904087
 Combine output records=89121
 Reduce input groups=67799
 Reduce shuffle bytes=1105621
 Reduce input records=89121
 Reduce output records=67799
 Spilled Records=178242
 Shuffled Maps =2
 Failed Shuffles=0
 Merged Map outputs=2
 GC time elapsed (ms)=566
 CPU time spent (ms)=20440
 Physical memory (bytes) snapshot=5465649152
 Virtual memory (bytes) snapshot=39886172160
 Total committed heap usage (bytes)=6074925056
Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
File Input Format Counters
 Bytes Read=5678681
File Output Format Counters
 Bytes Written=717253
```

With combiner

# Add combiner to the movie example

```
%%writefile codes/avgGenreCombiner.py
#!/usr/bin/env python

import sys
import csv
import json

genreList = {}

for line in sys.stdin:
 line = line.strip()
 genre, ratingString = line.split("\t", 1)
 ratingInfo = json.loads(ratingString)

 if genre in genreList:
 genreList[genre]["total_rating"] += ratingInfo["total_rating"]
 genreList[genre]["total_count"] += ratingInfo["total_count"]
 else:
 genreList[genre] = {}
 genreList[genre]["total_rating"] = ratingInfo["total_rating"]
 genreList[genre]["total_count"] = 1

for genre in genreList:
 print ("%s\t%s" % (genre, json.dumps(genreList[genre])))
```

- `hdfs dfs -rm -r intro-to-hadoop/output-movielens-06`
- `mapred streaming -input intro-to-hadoop/movielens/ratings.csv -output intro-to-hadoop/output-movielens-06 -file ./codes/avgGenreMapper02.py -mapper avgGenreMapper02.py -file ./codes/avgGenreReducer02.py -reducer avgGenreReducer02.py -file ./codes/avgGenreCombiner.py -combiner avgGenreCombiner.py -file ./movielens/movies.csv`
- How different are the number of shuffle bytes between the two jobs?

# Integrating Hadoop Job into Palmetto Workflow

```
#!/bin/bash

#PBS -N movieData
#PBS -l select=3:ncpus=8:mem=14gb
#PBS -l walltime=00:20:00
#PBS -j oe

cd ~/myhadoop

launch Hadoop cluster
./init_hadoop.sh

export environment variable
export HADOOP_CONF_DIR="/home/$USER/hadoop_palmetto/config/"

create directory and upload data
hdfs dfs -mkdir -p /user/$USER/intro-to-hadoop
hdfs dfs -put /zfs/citi/movielens intro-to-hadoop/

setup local data
mkdir movielens
hdfs dfs -get intro-to-hadoop/movielens/movies.csv movielens/

launch Hadoop job
mapred streaming -input intro-to-hadoop/movielens/ratings.csv -output intro-to-hadoop/output-movielens-02 -file ./codes/avgRatingMapper04.py -mapper avgRatingMapper04.py -file ./codes/avgRatingReducer01.py -reducer avgRatingReducer01.py -file ./movielens/movies.csv

get results
hdfs dfs -get intro-to-hadoop/output-movielens-02 .

stop Hadoop cluster
./stop_hadoop.sh
```

- With hdp module, Hadoop job can be invoked within a Palmetto PBS script, allowing the integration of large-data processing components into a standard HPC workflow.
  - Open a terminal, ssh to login001 (DUO required), and submit this script
    - `ssh login001`
    - `cd`
    - `cat -n ~/myhadoop/codes/movieAnalyzer.pbs`
    - `qsub ~/myhadoop/codes/movieAnalyzer.pbs`
  - View the final output when the job is finished
    - `qstat -anu $USER`
    - `cat part-00000 2>/dev/null | head -n 20`