

CPSC 4770/6770

## Distributed and Cluster Computing

Lecture 8: MPI Datatypes, Communicators and Groups

# MPI Datatypes

- MPI provides basic datatypes
  - MPI\_INT, MPI\_LONG, MPI\_CHAR, etc
- MPI also provides derived datatypes
  - Allows you to define your own data structures based upon sequences of the MPI primitive data types
    - Contiguous
    - Vector
    - Indexed
    - Struct
- New datatypes must be committed before they are used
  - int **MPI\_Type\_commit**(MPI\_Datatype \*datatype): Commits new datatype to the system
- Deallocate a specified datatype object
  - int MPI\_Type\_free(MPI\_Datatype \*datatype): Frees the datatype

Goal: Pack related data together  
to reduce total messages!

# Contiguous Derived Data Type

- int **MPI\_Type\_contiguous**(int count, MPI\_Datatype old\_type, MPI\_Datatype \*newtype): Produces a new data type by making *count* copies of an existing data type
- For example,

```
MPI_Type_contiguous(count, MPI_CHAR, &mytype);  
MPI_Type_commit(&mytype);  
MPI_Send(data, 1, mytype, 1, 99, MPI_COMM_WORLD);  
MPI_Type_free(&mytype);
```

is equivalent to:

```
MPI_Send(data, count, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
```

# contiguous.c

```
count = 4;  
MPI_Type_contiguous(count, MPI_FLOAT, &rowtype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
MPI_Send(&a[2][0], 1, rowtype, dest, tag, comm);
```

9.0	10.0	11.0	12.0
-----	------	------	------

1 element of  
rowtype

Create a data type representing a row of an array and distribute different rows to all processes

```
1 #include "mpi.h"  
2 #include <stdio.h>  
3 #define SIZE 4  
4  
5 main(int argc, char *argv[]) {  
6     int numtasks, rank, source=0, dest, tag=1, i;  
7     float a[SIZE][SIZE] =  
8         {1.0, 2.0, 3.0, 4.0,  
9          5.0, 6.0, 7.0, 8.0,  
10         9.0, 10.0, 11.0, 12.0,  
11         13.0, 14.0, 15.0, 16.0};  
12     float b[SIZE];  
13  
14     MPI_Status stat;  
15     MPI_Datatype rowtype; // required variable  
16  
17     MPI_Init(&argc,&argv);  
18     MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
19     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);  
20  
21     // create contiguous derived data type  
22     MPI_Type_contiguous(SIZE, MPI_FLOAT, &rowtype);  
23     MPI_Type_commit(&rowtype);  
24  
25     if (numtasks == SIZE) {  
26         // task 0 sends one element of rowtype to all tasks  
27         if (rank == 0) {  
28             for (i=0; i<numtasks; i++)  
29                 MPI_Send(&a[i][0], 1, rowtype, i, tag, MPI_COMM_WORLD);  
30         }  
31  
32         // all tasks receive rowtype data from task 0  
33         MPI_Recv(b, SIZE, MPI_FLOAT, source, tag, MPI_COMM_WORLD, &stat);  
34         printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f\n",  
35              rank, b[0], b[1], b[2], b[3]);  
36     }  
37     else  
38         printf("Must specify %d processors. Terminating.\n", SIZE);  
39  
40     // free datatype when done using it  
41     MPI_Type_free(&rowtype);  
42     MPI_Finalize();  
43 }
```

```
[[jin6@node0888 datatype]$ mpirun -np 4 contiguous  
rank= 0  b= 1.0 2.0 3.0 4.0  
rank= 1  b= 5.0 6.0 7.0 8.0  
rank= 2  b= 9.0 10.0 11.0 12.0  
rank= 3  b= 13.0 14.0 15.0 16.0
```

# Vector Derived Data Type

- `int MPI_Type_vector(int count, int blocklen, int stride, MPI_Datatype old_type, MPI_Datatype *newtype):` Creates a vector (strided) datatype
  - `count`: number of blocks
  - `blocklen`: number of elements in each block
  - `stride`: number of elements between start of each block
- It is similar to contiguous, but allows for regular gaps (stride) in the displacements

# vector.c

```
count = 4; blocklength = 1; stride = 4;  
MPI_Type_vector(count, blocklength, stride, MPI_FLOAT,  
                &columntype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
MPI_Send(&a[0][1], 1, columntype, dest, tag, comm);
```

2.0	6.0	10.0	14.0
-----	-----	------	------

1 element of  
columntype

Create a data type representing a column of an array and distribute different columns to all processes

```
1 #include "mpi.h"  
2 #include <stdio.h>  
3 #define SIZE 4  
4  
5 main(int argc, char *argv[]) {  
6     int numtasks, rank, source=0, dest, tag=1, i;  
7     float a[SIZE][SIZE] =  
8         {1.0, 2.0, 3.0, 4.0,  
9          5.0, 6.0, 7.0, 8.0,  
10         9.0, 10.0, 11.0, 12.0,  
11         13.0, 14.0, 15.0, 16.0};  
12     float b[SIZE];  
13  
14     MPI_Status stat;  
15     MPI_Datatype columntype; // required variable  
16  
17     MPI_Init(&argc, &argv);  
18     MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
19     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);  
20  
21     // create vector derived data type  
22     MPI_Type_vector(SIZE, 1, SIZE, MPI_FLOAT, &columntype);  
23     MPI_Type_commit(&columntype);  
24  
25     if (numtasks == SIZE) {  
26         // task 0 sends one element of columntype to all tasks  
27         if (rank == 0) {  
28             for (i=0; i<numtasks; i++)  
29                 MPI_Send(&a[0][i], 1, columntype, i, tag, MPI_COMM_WORLD);  
30         }  
31  
32         // all tasks receive columntype data from task 0  
33         MPI_Recv(b, SIZE, MPI_FLOAT, source, tag, MPI_COMM_WORLD, &stat);  
34         printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f\n",  
35                rank, b[0], b[1], b[2], b[3]);  
36     }  
37     else  
38         printf("Must specify %d processors. Terminating.\n", SIZE);  
39  
40     // free datatype when done using it  
41     MPI_Type_free(&columntype);  
42     MPI_Finalize();  
43 }  
44 }
```

```
[jin6@node0888 datatype]$ mpirun -np 4 vector  
rank= 0  b= 1.0 5.0 9.0 13.0  
rank= 2  b= 3.0 7.0 11.0 15.0  
rank= 3  b= 4.0 8.0 12.0 16.0  
rank= 1  b= 2.0 6.0 10.0 14.0
```


# Indexed Derived Data Type

- int **MPI\_Type\_indexed**(int count, int blocklens[], int indices[], MPI\_Datatype old\_type, MPI\_Datatype \*newtype): Creates an indexed datatype
  - count: number of blocks
  - blocklens: number of elements in each block
  - indices: displacement of each block in multiples of old\_type

# index.c


Create a datatype by extracting variable portions of an array and distribute to all tasks

count = 2;    blocklengths[0] = 4;  
              displacements[0] = 5;    blocklengths[1] = 2;  
                                      displacements[1] = 12;



`MPI_Type_indexed(count, blocklengths, displacements, MPI_FLOAT, &indextype);`

`MPI_Send(&a, 1, indextype, dest, tag, comm);`



1 element of  
indextype

```

1 #include "mpi.h"
2 #include <stdio.h>
3 #define NELEMENTS 6
4
5 main(int argc, char *argv[]) {
6     int numtasks, rank, source=0, dest, tag=1, i;
7     int blocklengths[2], displacements[2];
8     float a[16] =
9         {1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0,
10          9.0, 10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0};
11     float b[NELEMENTS];
12
13     MPI_Status stat;
14     MPI_Datatype indextype;    // required variable
15
16     MPI_Init(&argc,&argv);
17     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
19
20     blocklengths[0] = 4;
21     blocklengths[1] = 2;
22     displacements[0] = 5;
23     displacements[1] = 12;
24
25     // create indexed derived data type
26     MPI_Type_indexed(2, blocklengths, displacements, MPI_FLOAT, &indextype);
27     MPI_Type_commit(&indextype);
28
29     if (rank == 0) {
30         for (i=0; i<numtasks; i++)
31             // task 0 sends one element of indextype to all tasks
32             MPI_Send(a, 1, indextype, i, tag, MPI_COMM_WORLD);
33     }
34
35     // all tasks receive indextype data from task 0
36     MPI_Recv(b, NELEMENTS, MPI_FLOAT, source, tag, MPI_COMM_WORLD, &stat);
37     printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f %3.1f %3.1f\n",
38           rank, b[0], b[1], b[2], b[3], b[4], b[5]);
39
40     // free datatype when done using it
41     MPI_Type_free(&indextype);
42     MPI_Finalize();
43 }

```

```

[jin6@node0888 datatype]$ mpirun -np 4 index
rank= 2  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 3  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 0  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 1  b= 6.0 7.0 8.0 9.0 13.0 14.0

```



# Struct Derived Data Type


- int **MPI\_Type\_struct**(int count, int blocklens[], MPI\_Aint indices[], MPI\_Datatype old\_types[], MPI\_Datatype \*newtype): Creates a struct datatype
  - count: number of blocks
  - blocklens: number of elements in each block
  - indices: byte displacement of each block
  - old\_types: type of elements in each block
- int **MPI\_Type\_extent**( MPI\_Datatype datatype, MPI\_Aint \*extent): Returns the extent of a datatype

# MPI\_Type\_struct Example

```
typedef struct { float x,y,z,velocity; int n,type; } Particle;  
Particle particles[NELEM];
```

```
MPI_Type_extent(MPI_FLOAT, &extent);
```

```
count = 2; oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT  
offsets[0] = 0; offsets[1] = 4 * extent;  
blockcounts[0] = 4; blockcounts[1] = 2;
```



```
MPI_Type_struct(count, blockcounts, offsets, oldtypes, &particletype);
```

```
MPI_Send(particles, NELEM, particletype, dest, tag, comm);
```

Sends entire (NELEM) array of particles, each particle being comprised four floats and two integers.

# struct.c

Create a datatype by extracting variable portions of an array and distribute to all tasks

```
1 #include "mpi.h"
2 #include <stdio.h>
3 #define NELEM 25
4
5 main(int argc, char *argv[]) {
6     int numtasks, rank, source=0, dest, tag=1, i;
7
8     typedef struct {
9         float x, y, z;
10        float velocity;
11        int n, type;
12    } Particle;
13    Particle p[NELEM], particles[NELEM];
14    MPI_Datatype particletype, oldtypes[2]; // required variables
15    int blockcounts[2];
16
17    // MPI_Aint type used to be consistent with syntax of
18    // MPI_Type_extent routine
19    MPI_Aint offsets[2], extent;
20
21    MPI_Status stat;
22
23    MPI_Init(&argc,&argv);
24    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
25    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
26
27    // setup description of the 4 MPI_FLOAT fields x, y, z, velocity
28    offsets[0] = 0;
29    oldtypes[0] = MPI_FLOAT;
30    blockcounts[0] = 4;
31
32    // setup description of the 2 MPI_INT fields n, type
33    // need to first figure offset by getting size of MPI_FLOAT
34    MPI_Type_extent(MPI_FLOAT, &extent);
35    offsets[1] = 4 * extent;
36    oldtypes[1] = MPI_INT;
37    blockcounts[1] = 2;
38
```

```
39 // define structured type and commit it
40 MPI_Type_struct(2, blockcounts, offsets, oldtypes, &particletype);
41 MPI_Type_commit(&particletype);
42
43 // task 0 initializes the particle array and then sends it to each task
44 if (rank == 0) {
45     for (i=0; i<NELEM; i++) {
46         particles[i].x = i * 1.0;
47         particles[i].y = i * -1.0;
48         particles[i].z = i * 1.0;
49         particles[i].velocity = 0.25;
50         particles[i].n = i;
51         particles[i].type = i % 2;
52     }
53     for (i=0; i<numtasks; i++)
54         MPI_Send(particles, NELEM, particletype, i, tag, MPI_COMM_WORLD);
55 }
56
57 // all tasks receive particletype data
58 MPI_Recv(p, NELEM, particletype, source, tag, MPI_COMM_WORLD, &stat);
59
60 printf("rank= %d  %3.2f %3.2f %3.2f %3.2f %d %d\n", rank, p[3].x,
61        p[3].y, p[3].z, p[3].velocity, p[3].n, p[3].type);
62
63 // free datatype when done using it
64 MPI_Type_free(&particletype);
65 MPI_Finalize();
66 }
```

```
[jin6@node0888 datatype]$ mpirun -np 4 struct
rank= 2  3.00 -3.00 3.00 0.25 3 1
rank= 3  3.00 -3.00 3.00 0.25 3 1
rank= 0  3.00 -3.00 3.00 0.25 3 1
rank= 1  3.00 -3.00 3.00 0.25 3 1
```

# Overview of Communicators

- An MPI communicator (MPI\_Comm) is an object containing a group of processes that may communicate with each other
  - Predefined communicators
    - MPI\_COMM\_WORLD: contains all processes started with mpirun/mpiexec
    - MPI\_COMM\_SELF: contains just the local process itself, size is always 1
  - Create new communicators
    - Splitting the original communicator into n-parts
    - Creating subgroups of the original communicator
    - Re-ordering of processes based on topology information
    - Spawn new processes
    - Connect two applications and merge their communicators


# Splitting a Communicator

- `int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm *newcomm):` Creates new communicators based on colors and keys
  - All processes having the same color will be in the same subcommunicator
    - A process can just be part of one of the generated communicators
    - A process cannot “see” the other communicators
    - A process cannot “see” how many communicators have been created
    - If a process shall not be part of any of the resulting communicators
      - Set *color* to `MPI_UNDEFINED` (newcomm will be `MPI_COMM_NULL`)
  - Order processes with the same color according to the *key* value

`MPI_COMM_WORLD`

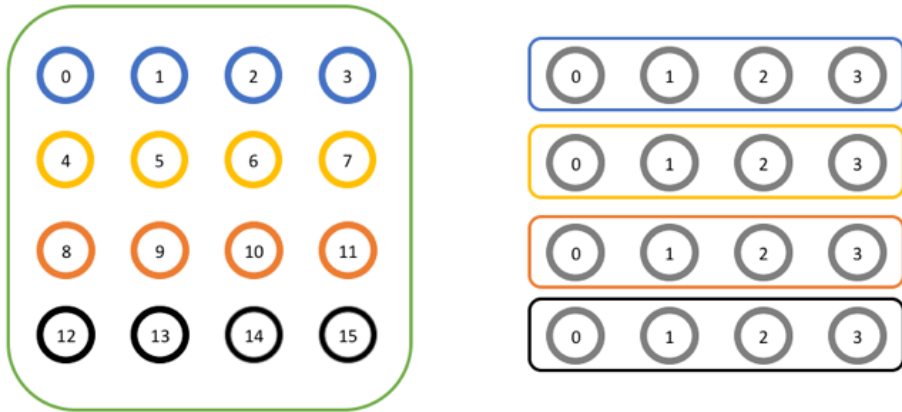
`newcomm, color=0, size = 4`

`newcomm, color=1, size = 3`



# MPI\_Comm\_split Example

Split a Large Communicator Into Smaller Communicators



WORLD RANK/SIZE: 0/16	ROW RANK/SIZE: 0/4
WORLD RANK/SIZE: 1/16	ROW RANK/SIZE: 1/4
WORLD RANK/SIZE: 2/16	ROW RANK/SIZE: 2/4
WORLD RANK/SIZE: 3/16	ROW RANK/SIZE: 3/4
WORLD RANK/SIZE: 4/16	ROW RANK/SIZE: 0/4
WORLD RANK/SIZE: 5/16	ROW RANK/SIZE: 1/4
WORLD RANK/SIZE: 6/16	ROW RANK/SIZE: 2/4
WORLD RANK/SIZE: 7/16	ROW RANK/SIZE: 3/4
WORLD RANK/SIZE: 8/16	ROW RANK/SIZE: 0/4
WORLD RANK/SIZE: 9/16	ROW RANK/SIZE: 1/4
WORLD RANK/SIZE: 10/16	ROW RANK/SIZE: 2/4
WORLD RANK/SIZE: 11/16	ROW RANK/SIZE: 3/4
WORLD RANK/SIZE: 12/16	ROW RANK/SIZE: 0/4
WORLD RANK/SIZE: 13/16	ROW RANK/SIZE: 1/4
WORLD RANK/SIZE: 14/16	ROW RANK/SIZE: 2/4
WORLD RANK/SIZE: 15/16	ROW RANK/SIZE: 3/4

```
// Get the rank and size in the original communicator
int world_rank, world_size;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

int color = world_rank / 4; // Determine color based on row

// Split the communicator based on the color and use the
// original rank for ordering
MPI_Comm row_comm;
MPI_Comm_split(MPI_COMM_WORLD, color, world_rank, &row_comm);

int row_rank, row_size;
MPI_Comm_rank(row_comm, &row_rank);
MPI_Comm_size(row_comm, &row_size);

printf("WORLD RANK/SIZE: %d/%d \t ROW RANK/SIZE: %d/%d\n",
       world_rank, world_size, row_rank, row_size);


MPI_Comm_free(&row_comm);
```

# Other Communicator Creation Functions

- int **MPI\_Comm\_dup**(MPI\_Comm comm, MPI\_Comm \*newcomm):  
Duplicates an existing communicator with all its cached information
- int **MPI\_Comm\_create**(MPI\_Comm comm, MPI\_Group group, MPI\_Comm \*newcomm): Creates a new communicator
  - group: a subset of the group of comm
  - Be aware of the difference between MPI\_Comm\_create\_group function (stay tuned)

# Overview of Groups

- An MPI Group (MPI\_Group) is the object describing the list of process forming a logical entity
  - A group has a size
    - MPI\_Group\_size
  - Every process in the group has a unique rank between 0 and (size of group-1)
    - MPI\_Group\_rank
  - A group is a local object, and can not be used for any communication





# Using MPI Groups

- `int MPI_Comm_group(MPI_Comm comm, MPI_Group *group):` Accesses the group associated with given communicator
- `int MPI_Group_union(MPI_Group group1, MPI_Group group2, MPI_Group *newgroup):` Produces a group by combining two groups
- `int MPI_Group_intersection(MPI_Group group1, MPI_Group group2, MPI_Group *newgroup):` Produces a group as the intersection of two existing groups
- `int MPI_Comm_create_group(MPI_Comm comm, MPI_Group group, int tag, MPI_Comm * newcomm):` Creates a new communicator
- `int MPI_Group_incl(MPI_Group group, int n, const int ranks[], MPI_Group *newgroup):` Produces a group by reordering an existing group and taking only listed members
- `int MPI_Group_excl(MPI_Group group, int n, const int ranks[], MPI_Group *newgroup):` Produces a group by reordering an existing group and taking only unlisted members

# incl.c

Creating a communicator which contains the prime ranks from MPI\_COMM\_WORLD







```
1 #include "mpi.h"
2 #include <stdlib.h>
3 #include <stdio.h>
4
5 int main(int argc, char *argv[]) {
6     // Get the rank and size in the original communicator
7     int world_rank, world_size;
8
9     MPI_Init(&argc, &argv);
10    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
11    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
12
13    // Get the group of processes in MPI_COMM_WORLD
14    MPI_Group world_group;
15    MPI_Comm_group(MPI_COMM_WORLD, &world_group);
16
17    int n = 7;
18    const int ranks[7] = {1, 2, 3, 5, 7, 11, 13};
19
20    // Construct a group containing all of the prime ranks in world_group
21    MPI_Group prime_group;
22    MPI_Group_incl(world_group, 7, ranks, &prime_group);
23
24    // Create a new communicator based on the group
25    MPI_Comm prime_comm;
26    MPI_Comm_create_group(MPI_COMM_WORLD, prime_group, 0, &prime_comm);
27
28    int prime_rank = -1, prime_size = -1;
29    // If this rank isn't in the new communicator, it will be
30    // MPI_COMM_NULL. Using MPI_COMM_NULL for MPI_Comm_rank or
31    // MPI_Comm_size is erroneous
32    if (MPI_COMM_NULL != prime_comm) {
33        MPI_Comm_rank(prime_comm, &prime_rank);
34        MPI_Comm_size(prime_comm, &prime_size);
35        printf("WORLD RANK/SIZE: %d/%d \t PRIME RANK/SIZE: %d/%d\n",
36              world_rank, world_size, prime_rank, prime_size);
37        MPI_Comm_free(&prime_comm);
38    }
39
40    MPI_Group_free(&world_group);
41    MPI_Group_free(&prime_group);
42
43    MPI_Finalize();
44 }
```

```
[jin6@node1684 08-datatype-communicator-group]$ mpirun -np 16 incl
WORLD RANK/SIZE: 1/16 PRIME RANK/SIZE: 0/7
WORLD RANK/SIZE: 2/16 PRIME RANK/SIZE: 1/7
WORLD RANK/SIZE: 3/16 PRIME RANK/SIZE: 2/7
WORLD RANK/SIZE: 5/16 PRIME RANK/SIZE: 3/7
WORLD RANK/SIZE: 7/16 PRIME RANK/SIZE: 4/7
WORLD RANK/SIZE: 11/16 PRIME RANK/SIZE: 5/7
WORLD RANK/SIZE: 13/16 PRIME RANK/SIZE: 6/7
```

Freeing groups and  
communicators

# group.c

MPI\_COMM\_WORLD



```

1 #include "mpi.h"
2 #include <stdlib.h>
3 #include <stdio.h>
4 #define NPROCS 10
5
6 main(int argc, char *argv[]) {
7     int rank, new_rank, sendbuf, recvbuf, numtasks,
8         ranks1[5]={0,1,2,3,4}, ranks2[5]={5,6,7,8,9};
9     MPI_Group orig_group, new_group; // required variables
10    MPI_Comm new_comm; // required variable
11
12    MPI_Init(&argc,&argv);
13    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
14    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
15
16    if (numtasks != NPROCS) {
17        printf("Must specify MP_PROCS= %d. Terminating.\n",NPROCS);
18        MPI_Finalize();
19        exit(0);
20    }
21
22    sendbuf = rank;
23
24    // extract the original group handle
25    MPI_Comm_group(MPI_COMM_WORLD, &orig_group);
26
27    // divide tasks into two distinct groups based upon rank
28    if (rank < NPROCS/2) {
29        MPI_Group_incl(orig_group, NPROCS/2, ranks1, &new_group);
30    } else {
31        MPI_Group_incl(orig_group, NPROCS/2, ranks2, &new_group);
32    }
33
34    // create new new communicator and then perform collective communications
35    MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);
36    MPI_Allreduce(&sendbuf, &recvbuf, 1, MPI_INT, MPI_SUM, new_comm);
37
38    // get rank in new group
39    MPI_Group_rank (new_group, &new_rank);
40    printf("rank= %d newrank= %d recvbuf= %d\n",rank,new_rank,recvbuf);
41
42    MPI_Finalize();
43 }

```

```

[jin6@node1555 group]$ mpicc group.c -o group
[jin6@node1555 group]$ mpirun -np 10 group
rank= 0 newrank= 0 recvbuf= 10
rank= 1 newrank= 1 recvbuf= 10
rank= 2 newrank= 2 recvbuf= 10
rank= 3 newrank= 3 recvbuf= 10
rank= 4 newrank= 4 recvbuf= 10
rank= 5 newrank= 0 recvbuf= 35
rank= 6 newrank= 1 recvbuf= 35
rank= 7 newrank= 2 recvbuf= 35
rank= 8 newrank= 3 recvbuf= 35
rank= 9 newrank= 4 recvbuf= 35

```