

CPSC 4770/6770

Distributed and Cluster Computing

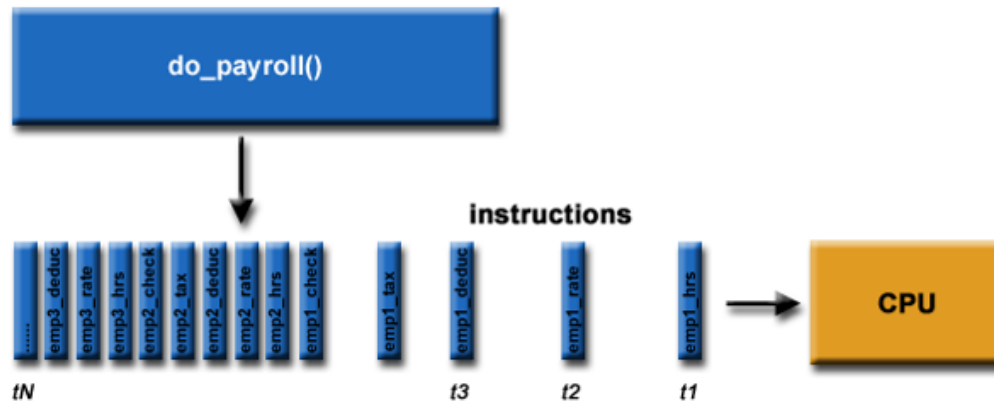
Lecture 1: Introduction to Parallel and Distributed Computing

New Palmetto Account Request


- Please fill out the account request from the link below to request a Palmetto account if you don't have one:
<https://www.palmetto.clemson.edu/palmetto/basic/new/>
- Make sure that you correctly select an Account Type (*Educational*), and an appropriate Rank (*Undergraduate* or *Graduate*)
- Your Sponsor's Clemson Email should be that of your course' instructor (jjin6@clemson.edu).
- ITHelp will approve the account once they received a correctly filled out form

What is Parallel Computing?

Sequential Computing



Parallel Computing



- The **problem** should be able to
 - Be broken apart into discrete pieces of work that can be solved simultaneously
 - Be solved in less time with multiple compute resources than with a single compute resource
- The **execution framework** should be able to
 - Execute multiple program instructions concurrently at any moment in time
- The **compute resource** might be
 - A single computer with multiple processors
 - An arbitrary number of computers connected by a network
 - A combination of both
 - A special computational component inside a single compute, separate from the main processors (GPU)

Source: https://computing.llnl.gov/tutorials/parallel_comp/


Measuring of Speed

- FLOPS: Floating-Point Operation Per Second
- MFLOPS = 1,000,000 FLOPS
- GFLOPS = 1,000,000,000 FLOPS
- TFLOPS = 1,000,000,000,000 FLOPS
 - Intel's ASCI Red was the first supercomputer in the world to achieve 1TFLOPS in 1993
- PFLOPS = 1,000,000,000,000,000 FLOPS
 - IBM RoadRunner was the first supercomputer to achieve 1 PFLOPS in 2008
 - US's Summit is the current fastest supercomputer in the world with a peak performance of 200 PFLOPS
- EFLOPS = 1,000,000,000,000,000,000 FLOPS

The Demand of Computation Speed

- Suppose whole global atmosphere divided into cells of size 1 mile x 1 mile x 1 mile to a height of 10 miles (10 cells high) – about 5×10^8 cells
- If each calculation requires 200 floating point operations, 10^{11} floating point operations necessary in one time step
- To forecast the weather for one week using 1-minute intervals, we need $60 \text{ (min/hr)} \times 24 \text{ (hr/day)} \times 7 \text{ (day)} = 10,080$ ($\sim 10^4$ time steps)
- The week-long forecast would require $10^{11} \times 10^4 = 10^{15}$ operations
- Maximum operations per clock tick for x86_64 CPU architecture: 4 (flop)
- A single-core CPU with 2.5GHz clock speed will have:
$$FLOPS = 2.5 \times 10^9 \times 4 = 10 \times 10^9 \text{ (flops)} = 10 \text{ (gflops)}$$
- Single computer simulation:
$$time = 10^{15} / (10 \times 10^9) / 3600 / 24 = 11.5 \text{ days}$$
- To do this within five minutes:
$$FLOPS = 10^{15} / 300 = 3.35 \times 10^{12} \text{ (flops)} = 3.4 \text{ (tflops)}$$
- PS4: theoretical peak performance of 1.84 (TFLOPS) using 8 cores

Computational Power for Weather Forecast Agencies (2013)




Source: <https://blogs.agu.org/wildwildscience/2013/02/17/seriously-behind-the-numerical-weather-prediction-gap/>

Catching up


- Two new supercomputers, Luna and Surge
- 2.89 PFLOPS each for a total of 5.78 PFLOPS (previous generation is only 776 TFLOPS)
- Increase water quantity forecast from 4000 locations to 2.7 million locations (700-fold increase in spatial density)
- Can track and forecast 8 storms at any given time
- 44.5 million dollars investment

Source: <https://www.noaa.gov/media-release/noaa-completes-weather-and-climate-supercomputer-upgrades>

Who is Using Parallel Computing?



Development Over Time



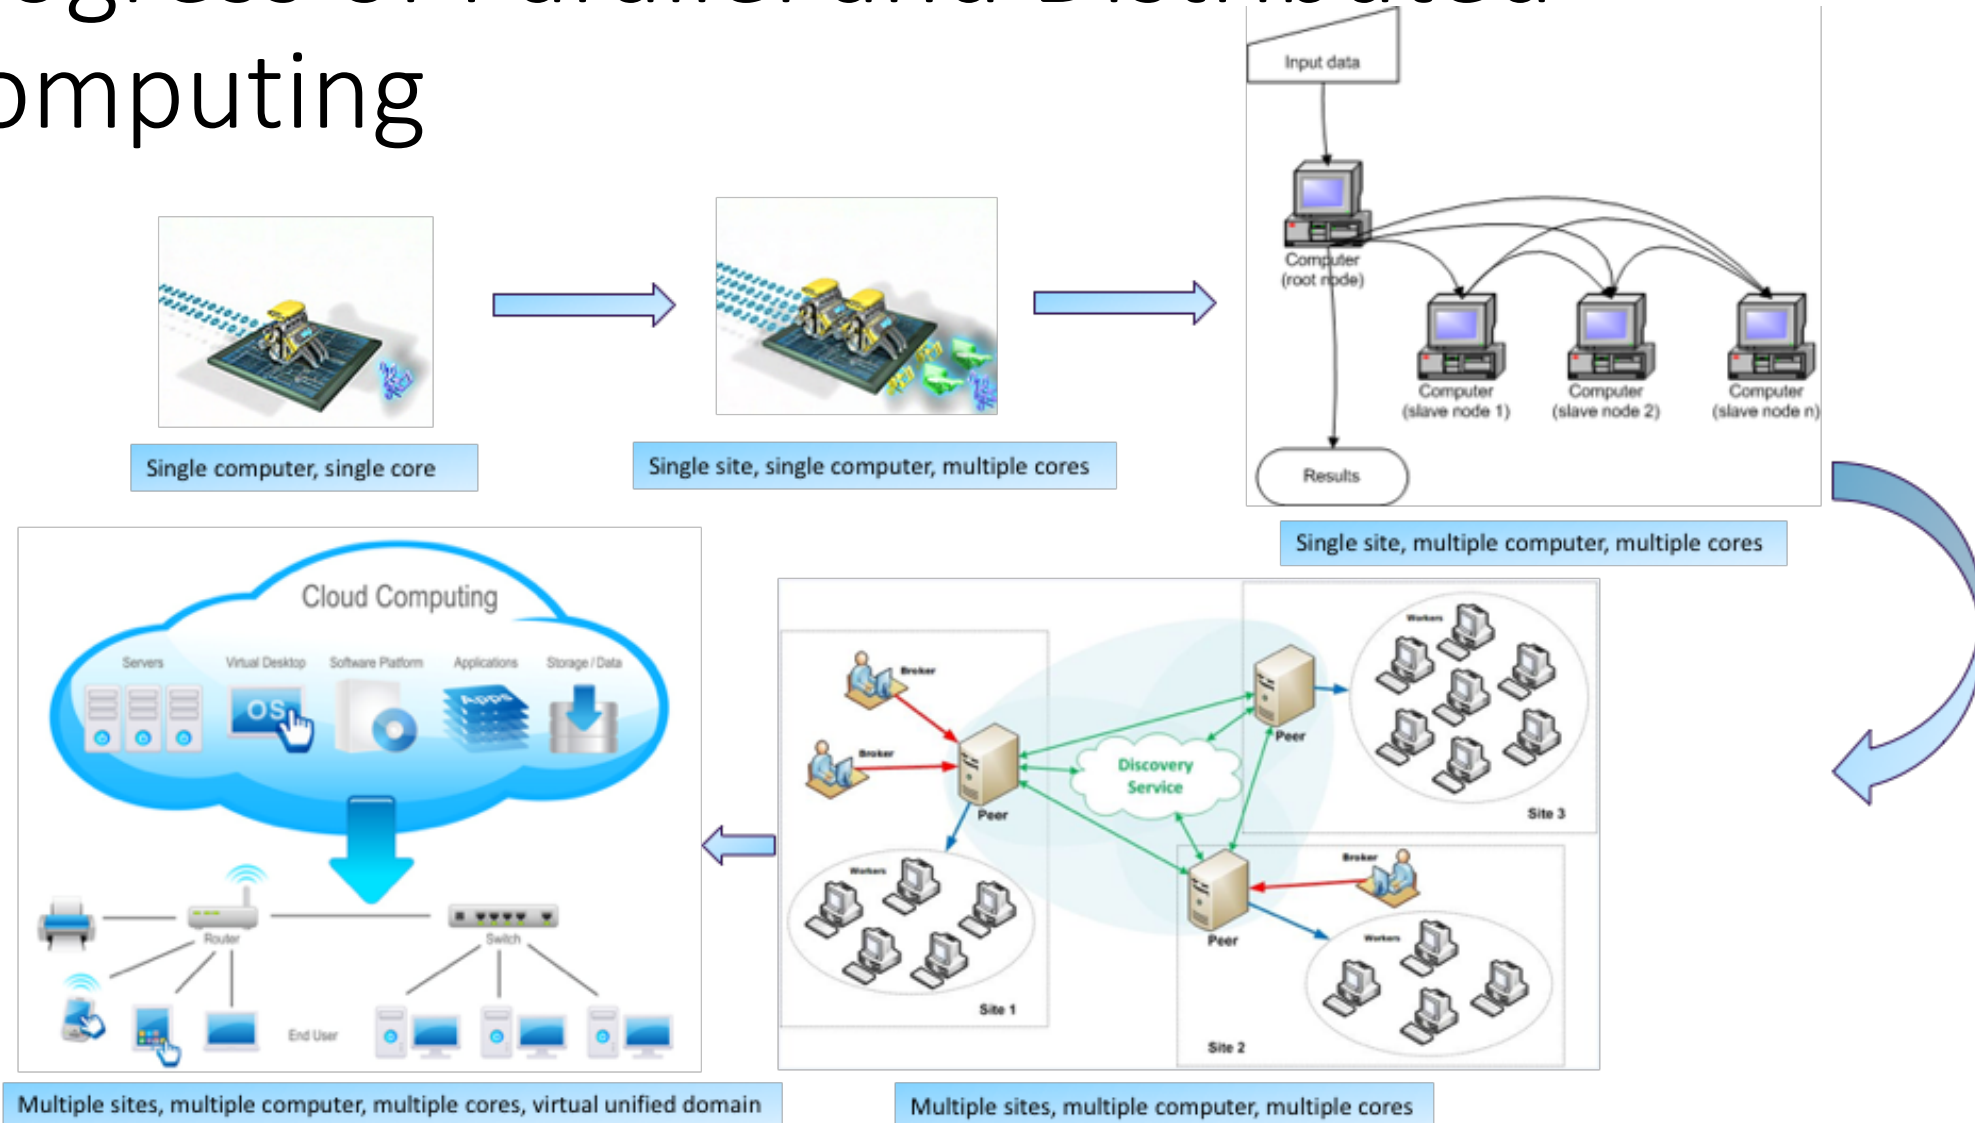
Source: <https://www.top500.org/statistics/perfdevel/>

Fraud Detection at PayPal

- 10M+ logins, 13M+ transactions, 300 variables per events
- ~4B inserts, ~8B selects
- MPI-like applications, Lustre Parallel File Systems, Hadoop
- Saved over \$700M in fraudulent transactions during their first year of deployment

Sources: <https://hpcuserforum.com/presentations/tuscon2013/IDCHPDABigDataHPC.pdf>
<https://www.intel.com/content/dam/www/public/us/en/documents/white-papers/big-data-meets-high-performance-computing-white-paper.pdf>

Progress of Parallel and Distributed Computing



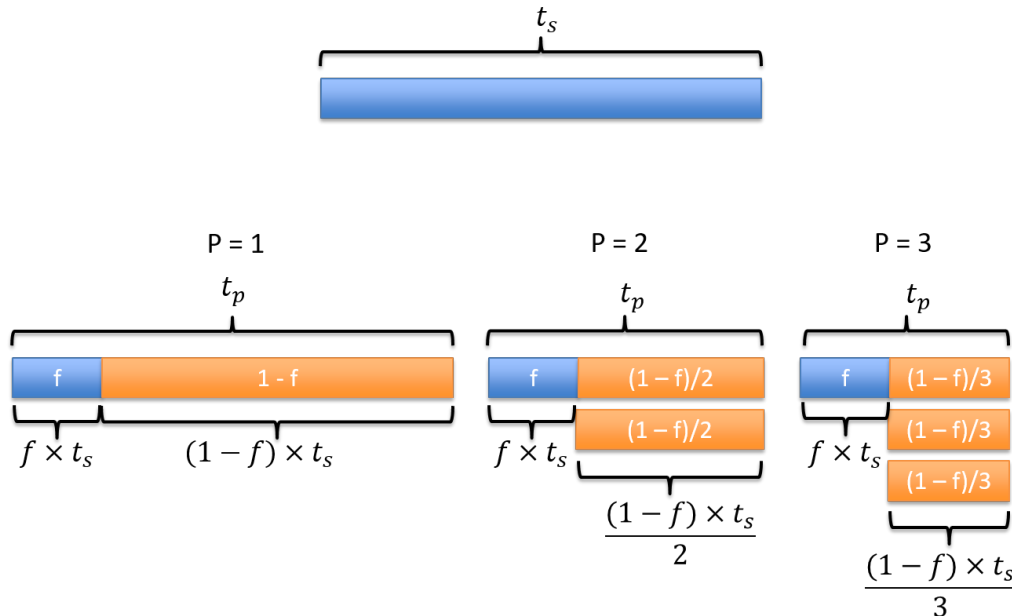
"Distributed Computing Systems are a collection of individual computing devices that **can communicate with each other.**" (Attiya and Welch, 2004)

Quantification of Performance Improvement

- **Parallel Speedup:** How much faster the program becomes once *some* computing resources are added

- $S(p) = \frac{\text{Sequential runtime}}{\text{Parallel runtime}} = \frac{t_s}{t_p}$

- Theoretical Max: Let f be the fraction of the program that is not parallelizable. Assume no communication overhead.



Amdahl's Law:

$$t_p = f t_s + \frac{(1-f)t_s}{p}$$
$$S(p) = \frac{t_s}{f t_s + \frac{(1-f)t_s}{p}} = \frac{p}{p f + 1 - f} = \frac{p}{(p-1)f + 1} \leq p$$

Superlinear speedup: $S(p) > p$ may occur when:

- Poor sequential reference implementation
- Memory caching
- I/O Blocking

Quantification of Performance Improvement

- **Parallel Efficiency:** Ratio of performance improvement per individual unit of computing resource

- $E = \frac{S(p)}{p} 100\% = \frac{1}{(p-1)f+1} 100\%$

- E.g., Suppose that 4% of my application is serial. What is my predicted speedup according to Amdahl's Law on 5 processors?

- $f = 0.04, p = 5$

- $E = \frac{1}{(p-1)*f+1} 100\% = \frac{1}{(5-1)*0.04+1} 100\% = 86.206897\%$

Limiting factors of $S(p)$:

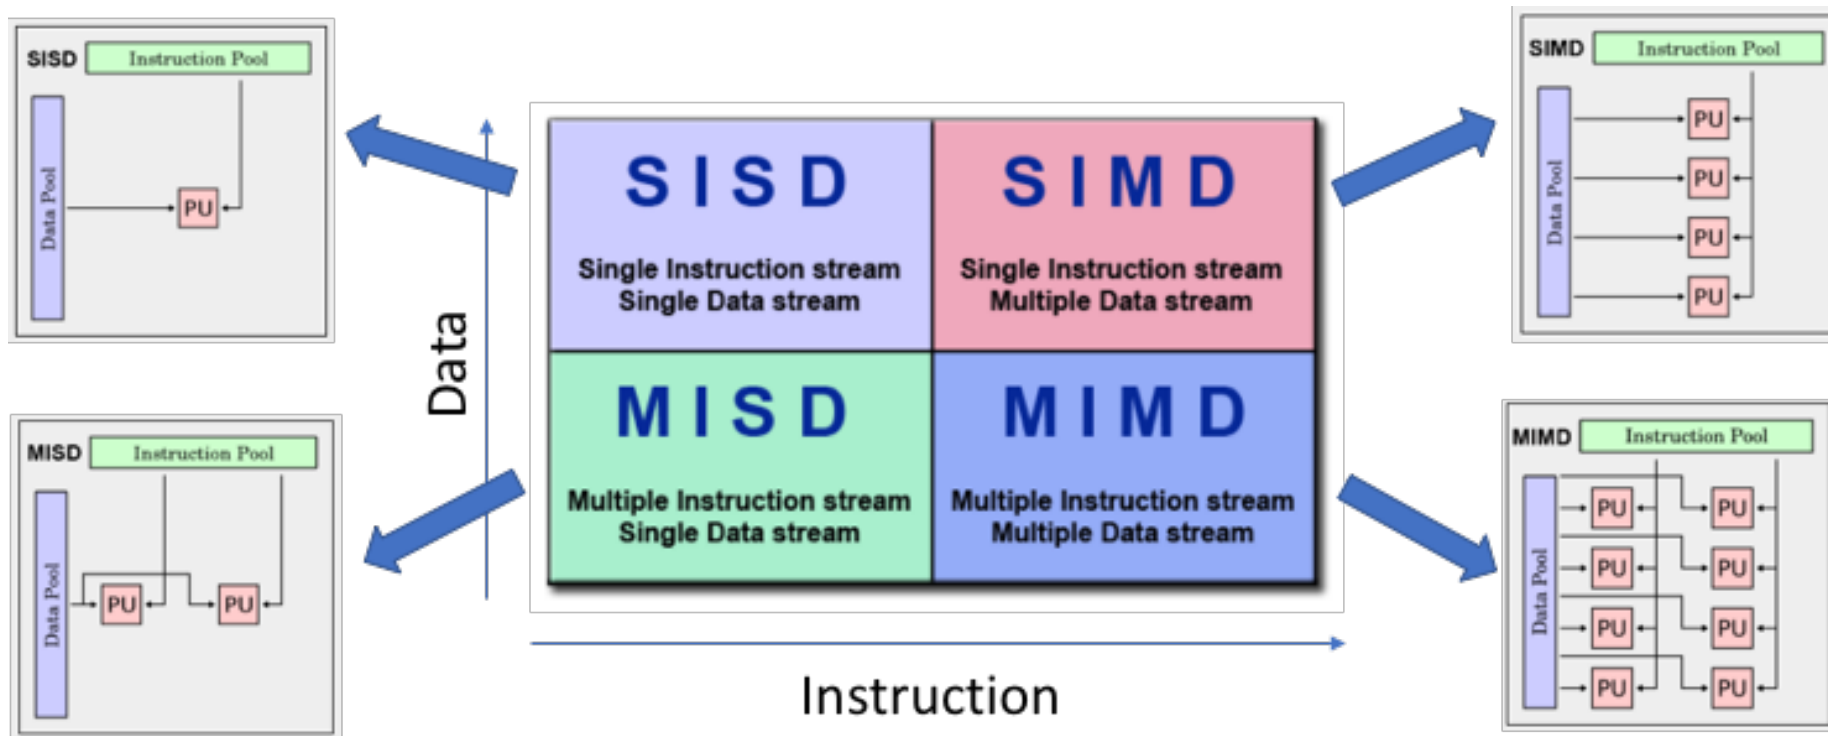
- Non-parallelization code
- Communication overhead

Exercise

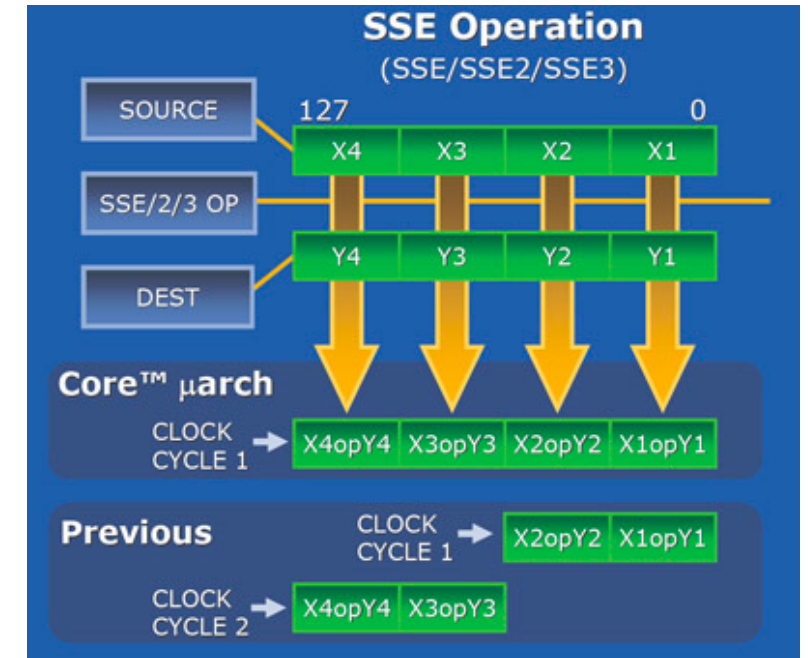
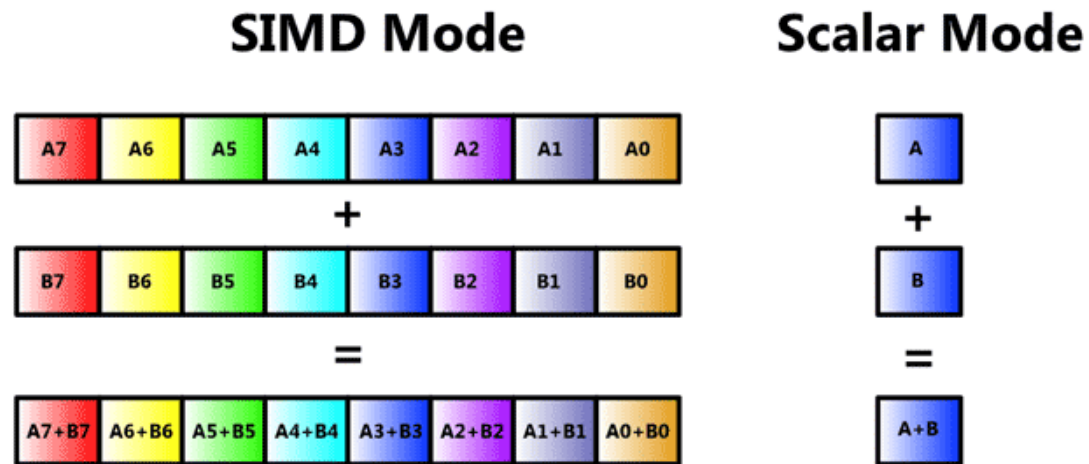
- Suppose that I get a speedup of 8 when I run my application on 10 processors.
 - According to Amdahl's Law, # What portion is serial?
 - What is the speedup on 20 processors? What is the efficiency?
 - What is the best speedup that I could hope for?

Types of Distributed Computing Systems

- Flynn's Taxonomy
 - Classifies multi-processor computer architectures into four different types according to how instructions and data flow through cores



Streaming SIMD




Source: <https://software.intel.com/en-us/articles/introduction-to-intel-advanced-vector-extensions>

- Intel's SSE
- AMD's 3DNow



Shared Memory

- One processor, multiple threads
- All threads have read/write access to the same memory
- Programming models:
 - Threads (pthread) – programmer manages all parallelism
 - OpenMP: Compiler extensions handle parallelization through in-code markers
 - Vendor libraries (e.g. Intel math kernel libraries)



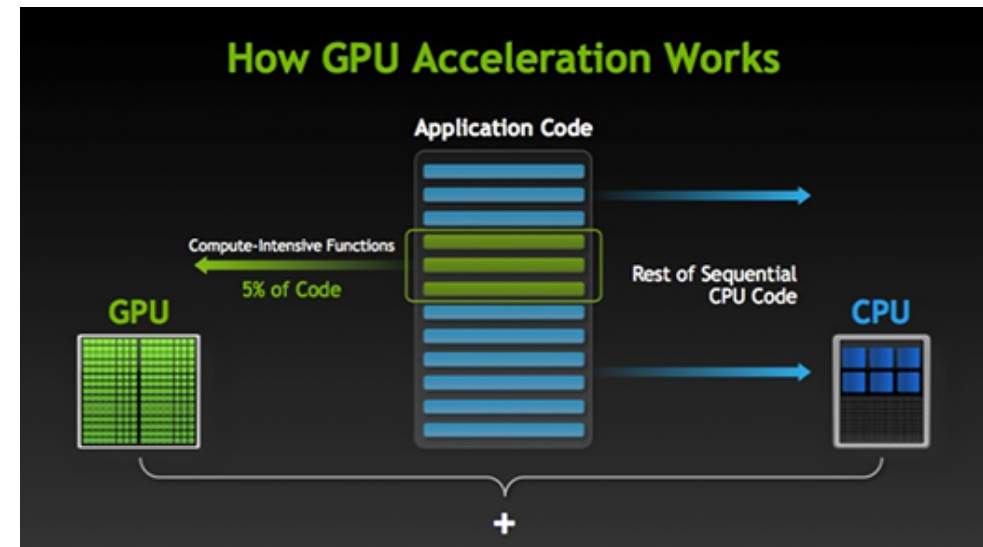
Message Passing

- Multiple processors, local memory
- Data passed via messages through communication network
- Programming models:
 - MPI: standardized message passing library
 - MPI + OpenMP (**hybrid model**)
 - MapReduce programming model



Heterogeneous Computing (Accelerators)

- GPU (Graphic Processing Units)
 - Processor unit on graphic cards designed to support graphic rendering (numerical manipulation)
 - Significant advantage for certain classes of scientific problem
 - CUDA – Library developed by NVIDIA for their GPUs
 - OpenACC – Standard developed by NVIDIA, Cray, and Portal Compiler (PGI)
 - OpenAMP – Extensions to Visual C++ (Microsoft) to direct computation to GPU
 - OpenCL – Set of standards by the group behind OpenGL
- FPGA (field programmable gate array)
 - Dynamically reconfigurable circuit board
 - Expensive, difficult to program
 - Power efficient, low heat



Benchmarking

- LINPACK (Linear Algebra Package): Dense Matrix Solver
- HPCC: High-Performance Computing Challenge
 - HPL (LINPACK to solve linear system of equation)
 - DGEMM (Double Precision General Matric Multiply)
 - STREAM (Memory bandwidth)
 - PTRANS (Parallel Matrix Transpose to measure processors communication)
 - RandomAccess (Random memory updates)
 - FFT (double precision complex discrete fourier transform)
 - Communication bandwidth and latency
- SHOC: Scalable Heterogeneous Computing - Non-traditional systems (GPU)
- TestDFSIO - I/O Performance of MapReduce/Hadoop Distributed File System

Ranking

- TOP500: Rank the supercomputers based on their LINPACK score
- GREEN500: Rank the supercomputers with emphasis on energy usage (LINPACK / power consumption)
- GRAPH500: Rank systems based on benchmarks designed for data-intensive computing